

Explaining Hopcroft, Tarjan, Gutwenger, and Mutzel's SPQR Decomposition Algorithm

SHOICHIRO YAMANISHI

yamanishi72@gmail.com

October 2, 2017

Abstract

The DFS based SPQR decomposition algorithm is explained. Role of TSTACK is explained. A simple case where the input graph does not have tree arc branches is explained as a sweeping algorithm. The simple case is expanded to the general case with tree arc branches and role of TSTACK with EOS is explained. Reason for the correction to ϕ calculation and the dynamic update of $HIGHPT(v)$ are given. Some implementation subtleties are explained.

I. OVERVIEW

A linear time SPQR decomposition algorithm using DFS was first presented by Hopcroft & Tarjan [HT73], and later it was corrected by Gutwenger & Mutzel [GM01]. The algorithm consists of 3 DFS explorations and some auxiliary subroutines. [GM01] isolates and clarifies the core of the algorithm called *check for type-2 pairs*, or denoted by `Type2()` and *check for type-1 pair*, or `Type1()`, as subroutines to the 3rd DFS called `PathSearch()`. Even after the clarification made by [GM01], it seems difficult to the author to understand the implementation of the algorithm presented in those papers. This article attempts to fill the gap between the literature and the implementation by explaining some key ideas with some examples. This article also clarifies some ambiguities on some parts of the algorithm that are not well explained in either of the papers above. It is strongly recommended the readers read [HT73] up to the proof of Lemma 17, and [GM01]. It is also recommended to read some other literature about the SPQR decomposition

to familiarise ourselves with basic concepts and terminology such as virtual edges and skeleton.

II. RESTRICTION TO SIMPLE GRAPHS

Here in this article only simple biconnected graphs are considered. Restricting the input graphs to be simple simplifies the DFS explorations and handling of split components when new components are found in `Type1()` and `Type2()` by eliminating the need to handle 2-cycles. This restriction is not severe, as we can handle multi-graphs by extra pre- and post-processing. We temporarily remove self-loops and bundle multi-edges into single edges before running the algorithm, and after the SPQR decomposition, the self-loops are restored, and additional P-nodes are further added for those multi-edges bundled.

With this restriction, the only places in the algorithm where multi-edges occur are in the split components of 3-bonds (minimal P-nodes) during `PathSearch()`, and in the split component of P-nodes during and after the merger

of consecutive 3-bonds into a maximal bonds in Algorithm 2 of [GM01]. The 3rd DFS, called `PathSearch()` is a destructive process where at each step up to 3 minimal split components can be separated from the input graph at a specific separation pair $\{a, b\}$, and the input graph denoted by G_c in [GM01] will be reduced and will become a new skeleton with a virtual edge $\{a, b\}$. The simplicity of the input graph will remain invariant through the process of `PathSearch()`, which makes the implementation simpler.

If a multi-graph without self-loops were given to the algorithm described as in [GM01], the separation pairs found in `Type1()` and `Type2()` will no longer be minimal. Also, a special care will have to be taken if the frond that ends a path is a multi-edge, in which case the lowest node (start of the path), and the highest node of the path form a separation pair candidate in `TSTACK` as the first element pushed, and then it could be falsely considered a separation pair when the lowest node of the path is visited in a post-traversal. With all those extra complications, the author thinks the best way to handle biconnected multi-graphs is to temporarily remove multi-edges and self-loops before running the algorithm.

III. ROLES OF VIRTUAL EDGES

A virtual edge has three roles.

- Specifies the location of the split pair in the split components.
- Keeps G_c biconnected. G_c is the graph explored by `PathSearch()`.
- Makes the skeleton tri-connected in the R nodes.

IV. KEY TO UNDERSTANDING THE ALGORITHM: TYPE 1 AND TYPE 2 SEPARATION PAIRS

The problem of finding the triconnected components and thus finding the SPQR-tree of a simple biconnected graph is reduced to finding the type-1 and type-2 pairs as **Lemma 3** in

[GM01] states (Fig. 1 and 2).




Figure 1: Type 1 Separation Pair




Figure 2: Type 2 Separation Pair

This is the key observation that leads to the algorithm found in [HT73]. We could develop a divide-and-conquer type algorithm at each split pair found, and applying it recursively to the (not necessarily minimal) split components after each split. The algorithm in [HT73] and [GM01] takes a clever iterative and recursive approach that runs in $O(|V| + |E|)$ time utilising DFS. It is iterative in a sense that the algorithm scans for a minimal split component along a path in the palm-tree in the descend-

ing order of node number assigned, and it is recursive on the paths found at the new frond or a tree branch.

Finding the minimal split components iteratively is achieved by a clever placement of `Type2()` and `Type1()` in `PathSearch()`. Finding them recursively is achieved by a clever use of node numbering and the EOS markers on `TSTACK` for each recursion.

In the following we first deal with the simple case without recursion and we study the core of the algorithm around the use of the three values on `TSTACK` for the type-2 pairs. Then we expand the discussion to the recursive case by explaining the role of the EOS markers and the modification of `TSTACK` in the transition to the new path in `PathSearch()`. Before walking through a simple case without recursion, we will briefly review `TSTACK` first.

V. ROLE OF TSTACK

We will review the meaning of h , a , & b , and some notations relevant to them with Fig. 2. Hereafter we will use node a to mean the node number as well as the identity of the node interchangeably. Node a and node b in an element of `TSTACK` are the lower and higher nodes of a type-2 separation pair candidate. Node b must be on the left most path from a child r_{j0} of node a . In [HT73] & [GM01], such a node like b is called the first descendant of r_{j0} , and it is denoted by $r_{j0} \xrightarrow{*} b$. In the same way, h is on the left most path from a child b_{i0} of b . The node h is also the highest descendant of b_{i0} . The value h is $b_{i0} + ND(b_{i0}) - 1$. The triplet (h, a, b) is enough to describe the corresponding minimal split components. The node induced by the split components is given by $D(b_{i0}) \cup \dots \cup D(b_n) \cup S \cup \{a, b\}$ according to the proof of **Lemma 17** in [HT73]. In the range of node numbers, it can be expressed as $[r_{j0}, h]$ and a . In the algorithm it can be simply expressed as $[a, h]$, and the condition to find the edges for the separation class is given by $a \leq x \leq h \wedge a \leq y \leq h$ in `Type2()`. This is because the nodes in $D(r_j), j < j_0$, i.e., r_j is a child of a to the left of r_j , have higher numbers than

h , and also from the way DFS is performed, and from the order in which the visited edge is pushed onto `ESTACK`, the nodes $r_j, j > j_0$, i.e., to the right of r_{j0} in the siblings of a , have yet been explored, and hence the edges to the right of S and $\{a, r_j\}$ are not yet pushed onto `ESTACK`. This is the reason why the range $a \leq x \leq h \wedge a \leq y \leq h$ works to identify the edges in the split component.

VI. DFS WALK 1: TREE WITHOUT BRANCHES

First, we deal with the case of a simple tree where tree arcs form a single straight path with no branches, which does not involve EOS on `TSTACK`, and in which h always coincides with b on `TSTACK`. This shows the basic idea of the sweeping algorithm of the type 2 part of `PathSearch()`. See Fig. 3.




Figure 3: G_c without tree branches

It is easy to see the following according to Algorithm 4: `PathSearch()` of [GM01].

- All the tree arcs from the root and the first frond $f1$ back to the root together form the first path.
- The second path onwards each consists of a single frond.

Whenever a frond that starts a path, (i.e., any frond other than $f1$) is visited, `TSTACK` is updated. If there is no element removed from `TSTACK`, which basically means if there is no frond criss-crossing the current frond, we will put (v, w, v) on top of the current stack (Fig. 4). $\{w, v\}$ is a separation pair candidate with

the corresponding separation class with edges incident to the vertices on the path from v to w . This will be tested when DFS exploration comes back to w , if it is still on TSTACK. On the other hand, if there are some elements removed from TSTACK, this means there are some fronds criss-crossing $\{w, v\}$ (Fig. 5). Because of the presence of v coming out of those $\{a_i, b_i\}, 0 \leq i \leq k$ pairs, those can no longer be candidates for a separation pair. A possibility is w as the lower node, and the highest (outer most) b_k . In this case, the corresponding split component will be specified by the highest (outer most) h_k , which coincides with b_k . Conceptually, those removed fronds on TSTACK and the frond $\{v, w\}$ are merged into a new frond as a split pair candidate on top of TSTACK. This is what happens in the bottom half of PathSearch() in [GM01], where a frond $(v \hookrightarrow w)$ is handled.




Figure 4: $v \hookrightarrow w$ not crossing others




Figure 5: $v \hookrightarrow w$ crossing others

One might wonder if there is a case where a_i are above or coincide with v , but we can show such a case can not happen. Suppose there is a pair $\{a_j, b_j\}$ such that a_j is above or coincides with v on top TSTACK. Then such a

pair would have been a proper separation pair found in Type2() when $(a_j \rightarrow r)$, where r is the child of a_j on the tree arc, was examined in PathSearch(), and such a pair would have been removed in it. A contradiction. Please note that if a_j coincides with v , then the tree arc $(v \rightarrow r)$ has been visited before the frond $(v \hookrightarrow w)$ from the way incidence list for each node is created by the ϕ value for each edge (Fig. 6). In general, when the DFS is leaving a node v , then TSTACK does not contain any element whose a is greater (higher) than v from the top down to EOS marker, which will be explained in the next section.




Figure 6: a_i & b_i evaluated before $(v \hookrightarrow w)$ is processed

There is one thing we have to take care after TSTACK is updated as above, which conceptually means merging the fronds into one. After the merger, the information about the arrivals of the fronds at a_i will be lost from TSTACK. This will cause a false finding of a separation pair as shown in Fig. 7. In Fig. 7, when the frond $(v \hookrightarrow w)$ is processed, two elements on TSTACK, (b_2, a_2, b_2) and (b_1, a_1, b_1) are merged to (b_1, w, b_1) , later when a_4 is being visited, it falsely finds (b_3, a_4, b_3) on TSTACK as a separation pair, but it is not due to the frond $(b_2 \hookrightarrow a_2)$, which has been lost from TSTACK.

To avoid finding such a false separation pair, we have to remove invalid fronds from TSTACK, and it is done at the end of a visit to $(v \rightarrow w)$ in PathSearch(). It removes elements from TSTACK whose h values are less than the current node v . For example when the DFS exploration is at the end of a visit to $\{a_1, r_1\}$ along the path, a_1 's high point is found to be b_1 . All the elements whose h value (and hence b value) is less than the high point are removed. This




Figure 7: Avoid false detection of separation pairs

prevents detection of false separation pairs. For the case in Fig. 7, at the end of the visit to a_2 , the high point of a_2 is found to be b_2 , and the element (b_3, a_4, b_3) is removed from TSTACK.

The algorithm proceeds in this way from the higher nodes down toward the tree root iteratively. This can be considered a sweeping algorithm. Along the exploration, at a node v , if there is an element on TSTACK whose a coincides with v , then this means there is no frond coming in or going out between a and b (not including a and b), and $\{a, b\}$ form a separation pair with all the edges between $\{a, b\}$ inclusive. See Fig. 8.




Figure 8: Order in which type-2 pairs are detected

Also we are sure that the split component found by (h, a, b) is minimal from the way the

fronds have been explored. If some elements on TSTACK share the same a , then the element with the closest b is always on top, and if some elements on TSTACK share the same b , then the element with the closest a is always on top.

After the separation class is removed, $\{a, b\}$ will be inserted to G_c as a virtual edge, the current top on TSTACK is removed, and $\text{Type2}()$ tries to find another separation pair with the same a ($=v$) if there is any on TSTACK. See 2, 3, 4 in Fig. 8. If not, then DFS exploration moves on to the next node below v . This sweeping continues until DFS exploration hits the root node, by which time all the separation pairs will have been found.

VII. HANDLING TYPE 1 CASES

The check for a type 1 separation pair is done after the checks for type 2 pairs are done, and the check is done only once. The reason for testing type 2 first, possibly multiple times, and then testing type 1 at most once, is as follows. At each current node v , G_c must be tested for Type 2 first to ensure the split component found is minimal, i.e., the split component should not be able to be split into two separation classes further by another separation pair in the node set induced by the separation class. Fig. 9 illustrates this. Suppose the current node for the DFS exploration is v , we see two separation pairs: $\{v, b\}$ as Type 2, and $\{v, \text{lowpt1}(v)\}$ as Type 1. If we test for Type 1 first, the separation class will contain another split pair $\{v, b\}$, which is against minimality. The minimality among the multiple split components found in $\text{Type2}()$ are minimal due to the way TSTACK is constructed, updated, and tested as shown above. Also, we don't have to test for Type 1 more than once. There can be another separation class with the same separation pair, but those will be found when DFS exploration comes back to v from a different neighbour (Fig. 10).




Figure 9: Detection of Type 1 pair




Figure 10: Two split components with the same Type 1 pair

VIII. HANDLING TRIANGLES (MINIMAL S-TYPE SEPARATION CLASS)

Triangles are handled as a special case in both `Type2()` and `Type1()`. `Type2()` deals with $v \rightarrow w \rightarrow x$, while `Type1` deals with $v \rightarrow w \hookrightarrow x$. The former can't be handled by `TSTACK`, and therefore it is detected by an explicit condition $\text{deg}(w) = 2 \wedge \text{firstChild}(w) > w$ in `Type2()` as shown in Fig. 11. The latter is detected as a minimal case for Type 1, which is $\text{lowpt1}(w) < v \wedge \text{lowpt2}(w) \geq v$ as shown in Fig. 12.




Figure 11: Type 2 triangle




Figure 12: Type 1 triangle

IX. ϕ VALUE AND ORDER OF INCIDENCE EDGES OF EACH NODE

[GM01] corrects the calculation of ϕ value of each edge as follows.

$$\phi(e) = \begin{cases} 3\text{lowpt1}(w) & \text{if } e = v \rightarrow w \wedge \text{lowpt2}(w) < v \\ 3w + 1 & \text{if } e = v \hookrightarrow w \\ 3\text{lowpt1}(w) + 2 & \text{if } e = v \rightarrow w \wedge \text{lowpt2}(w) \geq v \end{cases}$$

[GM01] states reason for the correction as to identify all multiple edges, but it is not clear why this correction is needed. Basically this is relevant to `Type1()` and the order in which the edges are pushed to `ESTACK`. Suppose `PathSearch()` is visiting v from w_{i0} in the post traversal as in Fig. 13, and it has detected a Type 1 separation pair $\{v, \text{lowpt1}(w_{i0})\}$, in which case $\text{lowpt2}(w_{i0}) \geq v$. Also suppose there is a frond $(v \hookrightarrow \text{lowpt1}(w_{i0}))$. We must have visited it before visiting the tree arc $(v \rightarrow w_{i0})$, so that the it is the first edge pushed to `ESTACK` and hence the last one popped in the following if clause after the while loop in `Type1()`.

`ESTACK.top() = (v, lowpt1(w)) then`

The ϕ calculation proposed in [GM01] ensures this order, but the one proposed in [HT73] does not.




Figure 13: An edge between a type-1 separation pair

X. DFS WALK 2: GENERAL CASE WITH TREE BRANCHES AND EOS

Finally, we will expand the discussion to the general case for the palm-trees with tree branches. First, imagine there is a branch from a node v as shown in Fig. 14. Here, the tree arc $(v \rightarrow w)$ marks the start of a new path that follows the left-most tree arcs and ends with the deepest frond, whose destination node corresponds to $lowpt1(w)$. TSTACK is updated according to `PathSearch()`. The basic idea is the same as for the single path case, and the path $v \rightarrow w \xrightarrow{*} h \leftrightarrow lowpt1(w)$ can be considered a frond from v to $lowpt1(w)$. The criss-crossing fronds with the current ones are removed from TSTACK, merged with the current one, and a single new element is pushed to TSTACK. In Fig. 14, the two elements (b_1, a_1, b_1) and (b_2, a_2, b_2) are removed from TSTACK and a new element $(b_2, lowpt1(w), b_2)$ is pushed. The difference is that now h is different from b . Node b is at the branching node on the current path, which is v , and node h is the highest possible node along the left-most path from node w . Node a is lowest reachable from w , which is $lowpt1(w)$. Then the modification of TSTACK for the current path is done, and an EOS marker is pushed on to TSTACK to prepare for the recursive processing for the new path that starts with $(v \rightarrow w)$.

Here we observe the following two invariants on the current path.

- The conceptual (merged) fronds formed by $(a \leftrightarrow b)$ on TSTACK are not criss crossing.
- The nodes for b are always on the current path.




Figure 14: New path from v

We also observe the nodes for a on TSTACK can no longer be on the current path. **Lemma 3** in [GM01] requires b to be on the left-most path from a child of node a . The following explains how the algorithm handles this condition properly. See Fig. 15. In this case the current path starts from node c along the thicker line. At some point at a tree branch, TSTACK gets an element $(h, v, lowpt1(w))$, and $lowpt1(w)$ is located down on an earlier path. If the DFS visits c in the post traversal, it checks if c coincides with node a on top of TSTACK. If it does (Fig. 15 b), then the top of TSTACK will be a separation pair. If it does not (Fig. 15 a), then all the elements for the current paths including EOS marker are removed from TSTACK at the end of the visit.

To avoid false detection of a separation pair, `PathSearch()` removes elements from TSTACK at the end of a visit to a tree arc just like the case of a simple path described before. All the elements whose h is less than $HIGHT(v)$ are removed. The reason why h is used instead of

b is that unlike for the case of a simple path, b and h can have different values, and there can be multiple elements with different h value but with the same b value. Fig. 16 is such a case. This is equivalent to avoiding finding false separation pair to the left of $D(b_{i0})$ in Fig. 2.

Overall, whenever the DFS enters a new tree arc that starts a path, update the elements on TSTACK as shown in PathSearch() in [GM01], place EOS on top to temporarily suspend the processing on the current path, then move on to the newly discovered path. On the new path the DFS tries to find separation pairs along the new path, and also when the DFS comes back to the node where the new path has started, e.g., c in Fig. 15 (b), if there is an element on TSTACK whose a coincides with the node, then the separation pair and the split components are processed. This is done recursively whenever a new path is found along the DFS exploration.




Figure 15: a not necessarily on the current path

There is a subtle point in updating the TSTACK in the beginning of the tree arc that starts a path. If some elements are removed from TSTACK, then the outer-most b will be used for the new element to be pushed. The outer-most b coincides with b in the last element removed. However, we can't simply use h in the last element. The reason is that the elements from fronds and elements from tree arcs that start new paths that share the same b

are intertwined on TSTACK. See Fig. 17. In this case the element from the frond ($v \hookrightarrow w_1$) is at the bottom. Later when the frond ($x \hookrightarrow y$) is processed those three elements are removed from TSTACK and a new h is not from the last element removed due to the presence of the element from fronds.

One might wonder if a similar treatment is necessary when elements of TSTACK whose h are less than $\text{HIGHPT}(v)$ are removed. For example if there is a frond in Fig. 16 such that $(b \hookrightarrow a_0), a_0 < a_1$. In this case there can be (b, a_0, b) below (h_1, a_1, b) on TSTACK when the DFS is visiting v , and the removal stops when it sees (h_1, a_1, b) , and (b, a_0, b) will go unremoved. However this might cause a false detection of separation pair. If (b, a_0, b) is still on TSTACK when the DFS is visiting a_1 , then $\text{HIGHPT}(a_1) > b_1$, and (b, a_0, b) will be removed there before the DFS reaches a_0 in the post traversal.




Figure 16: Removing invalid elements from TSTACK

XI. REASON WHY DYNAMIC UPDATES ARE NEEDED FOR HIGHPT

[GM01] states the need for updates of HIGHPT for each node on-the-fly, but it does not state the reason. Fig. 18 explains this in relation to invalid elements removal on TSTACK at the end of each visit to v explained above. In Fig.




Figure 17: Order of elements that share the same b on TSTACK

REFERENCES

- [HT73] Hopcroft, J.E. and Tarjan, R.E. (1973).
Dividing a graph into triconnected components *SIAM J. Comput.*, 2(3):135-158.
- [GM01] Gutwenger, C. and Mutzel P. (2001).
A linear time implementation of SPQR-trees *Proc. 8th International Symposium on Graph Drawing (GD2000)*, Lecture Notes in Computer Science 1984, Springer-Verlag, pp. 77-90

18 (a), node a' is the current node being visited, and it finds (h_1, a_1, b_1) on top of the stack. Supposed $\text{HIGHPT}(a') = b'$. In Fig. 18 (a), the subgraph from b_1 - b' - h' still remains in G_c and hence (h_1, a_1, b_1) has to be removed. However, there are some cases in which by the time DFS visits node a' on its way back (post-traversal), the subgraph induced by b_1 , b' and h' has been removed from G_c as separation classes. In this case (h_1, a_1, b_1) is a valid separation pair candidate and it should be kept on TSTACK. Since whether or not to remove (h_1, a_1, b_1) is determined by $\text{HIGHPT}(a')$, keeping it up-to-date during the DFS traversal is necessary.




Figure 18: Need to update HIGHPT